

Chapter 6


Fitting: Voting and Hough Transform

Prof. Fei-Fei Li, Stanford University

Contents

Line fitting

- Hough Transform
- RANSAC (**RAN**dom **SA**mples **C**onsensus)





Fitting as Search in Parametric Space

- *Choose a parametric model to represent a set of features*
- *Membership criterion is not local*
 - *Can't tell whether a point belongs to a given model just by looking at that point.*
- *Three main questions:*
 - *What model represents this set of features best?*
 - *Which of several model instances gets which feature?*
 - *How many model instances are there?*
- *Computational complexity is important*
 - *It is infeasible to examine every possible set of parameters and every possible combination of features*




Example: Line Fitting

- Why fit lines? Many objects characterized by presence of straight lines*



- Wait, why aren't we done just by running edge detection?*

Difficulty of Line Fitting :



- **Extra edge points (clutter), multiple models:**
 - Which points go with which line, if any?
- **Only some parts of each line detected, and some parts are missing:**
 - How to find a line that bridges missing evidence?
- **Noise in measured edge points, orientations:**
 - How to detect true underlying parameters?



Voting

- *It's not feasible to check all combinations of features by fitting a model to each possible subset.*
- *Voting is a general technique where we let the features vote for all models that are compatible with it.*
 - *Cycle through features, cast votes for model parameters.*
 - *Look for model parameters that receive a lot of votes.*
- *Noise & clutter features will cast votes too, but typically their votes should be inconsistent with the majority of “good” features.*
- *Ok if some features not observed, as model can span multiple fragments.*


*Slide credit: Kristen
Grauman*

Fitting Lines

- *Given points that belong to a line, what is the line?*
- *How many lines are there?*
- *Which points belong to which lines?*
- **Hough Transform** is a voting technique that can be used to answer all of these
- *Main idea:*
 1. *Record all possible lines on which each edge point lies.*
 2. *Look for lines that get many votes.*




Finding Lines in an Image: Hough Space



- **Connection between image (x,y) and Hough (m,b) spaces**
 - A line in the image corresponds to a point in Hough space.
 - To go from image space to Hough space:
 - Given a set of points (x,y) , find all (m,b) such that $y = mx + b$

Finding Lines in an Image: Hough Space



- **Connection between image (x,y) and Hough (m,b) spaces**
 - A line in the image corresponds to a point in Hough space.
 - To go from image space to Hough space:
 - Given a set of points (x,y) , find all (m,b) such that $y = mx + b$
 - What does a point (x_0, y_0) in the image space map to?
 - Answer: the solutions of $b = -x_0m + y_0$
 - This is a line in Hough space

Finding Lines in an Image: Hough Space





Image space



Hough (parameter) space

- ***What are the line parameters for the line that contains both (x_0, y_0) and (x_1, y_1) ?***
 - ***It is the intersection of the lines $b = -x_0m + y_0$ and $b = -x_1m + y_1$***

Finding Lines in an Image: Hough Space





Image space




Hough (parameter) space

- ***How can we use this to find the most likely parameters (m,b) for the most prominent line in the image space?***
 - ***Let each edge point in image space vote for a set of possible parameters in Hough space***
 - ***Accumulate votes in discrete set of bins; parameters with the most votes indicate line in image space.***

Polar Representation for Lines

- *Issues with usual (m,b) parameter space:*

- *Can take on infinite values;*
- *Undefined for vertical lines.*



d : perpendicular distance from line to origin

θ : angle the perpendicular line makes with the x-axis

$$x \cos \theta + y \sin \theta = d$$

where $\theta \in [0, \pi)$ and $d \in \mathbb{R}$

- *Point in image space \Rightarrow sinusoid segment in Hough space*

Hough Transform Algorithm

Using the polar parameterization:


$$x \cos \theta + y \sin \theta = d$$

Basic Hough transform algorithm

1. *Initialize $H[d, \theta] = 0$.*
2. *For each edge point (x, y) in the image for $\theta \in [0, \pi)$ $H[d, \theta] += 1$*
3. *Find the value(s) of (d, θ) where $H[d, \theta]$ is maximum*
4. *The detected line in the image is given by*

$$d = x \cos \theta + y \sin \theta$$

- *Time complexity (in terms of number of votes)?*



Example: HT for Straight Lines





Image space edge coordinates



θ


Votes

Bright value = high vote count



Black = no votes

Example: HT for Straight Lines




Square:





Example: HT for Straight Lines




Real-World Examples






Showing longest segments found

Impact of Noise on Hough Transform




*Image space
edge coordinates*




Votes

What difficulty does this present for an implementation?

Impact of Noise on Hough Transform



*Image space
edge coordinates*



Votes

Here, everything appears to be “noise”, or random edge points, but we still see peaks in the vote space.

Voting: Practical Tips

- *Minimize irrelevant tokens first (take edge points with significant gradient magnitude)*
- *Choose a good grid / discretization*
 - *Too coarse: large votes obtained when too many different lines correspond to a single bucket*
 - *Too fine: miss lines because some points that are not exactly collinear cast votes for different buckets*
- *Vote for neighbors, also (smoothing in accumulator array)*
- *Utilize direction of edge to reduce free parameters by 1*
- *To read back which points voted for “winning” peaks, keep tags on the votes.*

Hough Transform: Pros and Cons

Pros

- *All points are processed independently, so can cope with occlusion*
- *Some robustness to noise: noise points unlikely to contribute consistently to any single bin*
- *Can detect multiple instances of a model in a single pass*

Cons

- *Complexity of search time increases exponentially with the number of model parameters*
- *Non-target shapes can produce spurious peaks in parameter space*
- *Quantization: hard to pick a good grid size*

Another model fitting strategy: RANSAC

[Fischler & Bolles 1981]

- ***RAN***dom ***SA***mples ***C***onsensus
- ***Approach: we want to avoid the impact of outliers, so let's look for "inliers", and use only those.***
- ***Intuition: if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.***


RANSAC

RANSAC loop:

1. *Randomly select a seed group of points on which to base transformation estimate (e.g., a group of matches)*
 2. *Compute transformation from seed group*
 3. *Find inliers to this transformation*
 4. *If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers*
- *Keep the transformation with the largest number of inliers*


RANSAC Line Fitting Example

- ***Task: Estimate the best line***
 - *How many points do we need to estimate the line?*




RANSAC Line Fitting Example

- *Task: Estimate the best line*




RANSAC Line Fitting Example

- *Task: Estimate the best line*



RANSAC Line Fitting Example


- *Task: Estimate the best line*



*Total number of
points within a
threshold of line.*


RANSAC Line Fitting Example

- *Task: Estimate the best line*



RANSAC Line Fitting Example


- *Task: Estimate the best line*



Repeat, until we get a good result.

RANSAC Line Fitting Example

- *Task: Estimate the best line*



Algorithm 15.4: RANSAC: fitting lines using random sample consensus

Determine:

n — the smallest number of points required

k — the number of iterations required

t — the threshold used to identify a point that fits well

d — the number of nearby points required
to assert a model fits well

Until k iterations have occurred

Draw a sample of n points from the data
uniformly and at random

Fit to that set of n points

For each data point outside the sample

Test the distance from the point to the line
against t ; if the distance from the point to the line
is less than t , the point is close

end

If there are d or more points close to the line
then there is a good fit. Refit the line using all
these points.

end

Use the best fit from this collection, using the
fitting error as a criterion

RANSAC: How many samples?

- *How many samples are needed?*
 - *Suppose w is fraction of inliers (points from line).*
 - *n points needed to define hypothesis (2 for lines)*
 - *k samples chosen.*
- *Prob. that a single sample of n points is correct: w^n*
- *Prob. that all k samples fail is: $(1 - w^n)^k$*


Choose k high enough to keep this below desired failure rate.

RANSAC: Computed k ($p=0.99$)

| Sample size n | 5% | 10% | Proportion of outliers | | | 40% | 50% |
|------------------|----|-----|------------------------|-----|-----|-----|------|
| | | | 20% | 25% | 30% | | |
| 2 | 2 | 3 | 5 | 6 | 7 | 11 | 17 |
| 3 | 3 | 4 | 7 | 9 | 11 | 19 | 35 |
| 4 | 3 | 5 | 9 | 13 | 17 | 34 | 72 |
| 5 | 4 | 6 | 12 | 17 | 26 | 57 | 146 |
| 6 | 4 | 7 | 16 | 24 | 37 | 97 | 293 |
| 7 | 4 | 8 | 20 | 33 | 54 | 163 | 588 |
| 8 | 5 | 9 | 26 | 44 | 78 | 272 | 1177 |

After RANSAC

- *RANSAC divides data into inliers and outliers and yields estimate computed from minimal set of inliers.*
- *Improve this initial estimate with estimation over all inliers (e.g. with standard least-squares minimization).*
- *But this may change inliers, so alternate fitting with re-classification as inlier/outlier.*




RANSAC: Pros and Cons

- **Pros:**
 - *General method suited for a wide range of model fitting problems*
 - *Easy to implement and easy to calculate its failure rate*
- **Cons:**
 - *Only handles a moderate percentage of outliers without cost blowing up*
 - *Many real problems have high rate of outliers (but sometimes selective choice of random subsets can help)*
- *The Hough transform can handle high percentage of outliers*

Generalized Hough Transform

- *What if want to detect arbitrary shapes defined by boundary points and a reference point?*



*At each boundary point,
compute displacement vector:*

$$\mathbf{r} = \mathbf{a} - \mathbf{p}_i.$$

*For a given model shape: store
these vectors in a table indexed
by gradient orientation θ .*

Dana H. Ballard, [Generalizing the Hough Transform to Detect Arbitrary Shapes](#), 1980


Generalized Hough Transform

To detect the model shape in a new image:


- *For each edge point*
 - *Index into table with its gradient orientation θ*
 - *Use retrieved r vectors to vote for position of reference point*
- *Peak in this Hough space is reference point with most supporting edges*

Assuming translation is the only transformation here, i.e., orientation and scale are fixed.


Example: Generalized Hough Transform



Example: Generalized Hough Transform




Example: Generalized Hough Transform




*Range of voting
locations for test point*

Example: Generalized Hough Transform




***Range of voting
locations for test point***

Example: Generalized Hough Transform




Example: Generalized Hough Transform




*Displacement vectors
for model points*

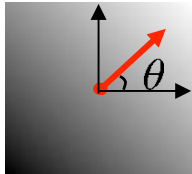
Example: Generalized Hough Transform



Example: Generalized Hough Transform



Extensions


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Extension 1: Use the image gradient

1. *same*
2. *for each edge point $I[x,y]$ in the image*

$\theta = \text{gradient at } (x,y)$

$d = x \cos \theta + y \sin \theta$

$H[d, \theta] += 1$

3. *same*

4. *same*

(Reduces degrees of freedom)

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

Extensions

Extension 2

- *Give more votes for stronger edges (use magnitude of gradient)*

Extension 3

- *Change the sampling of (d, θ) to give more/less resolution*

Extension 4

- *The same procedure can be used with circles, squares, or any other shape...*


Extension: Cascaded Hough Transform

- Let's go back to the original (m, b) parametrization
- A line in the image maps to a pencil of lines in the Hough space
- What do we get with parallel lines or a pencil of lines?
 - Collinear peaks in the Hough space!
- So we can apply a Hough transform to the output of the first Hough transform to find vanishing points
- T. Tuytelaars, M. Proesmans, L. Van Gool ["The cascaded Hough transform"](#), *ICIP'97*.

Finding Vanishing Points




Cascaded Hough Transform



T. Tuytelaars, M. Proesmans, L. Van Gool ["The cascaded Hough transform"](#), ICIIP'97.

Cascaded Hough Transform




T. Tuytelaars, M. Proesmans, L. Van Gool "[The cascaded Hough transform](#)", ICIP'97.

Hough Transform for Circles

- *Circle: center (a,b) and radius r*

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- *For a fixed radius r , unknown gradient*




Hough Transform for Circles

- **Circle: center (a,b) and radius r**

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- **For a fixed radius r , unknown gradient direction**




Hough Transform for Circles

- **Circle: center (a,b) and radius r**

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- **For an unknown radius r , unknown gradient direction**



Hough Transform for Circles

- **Circle: center (a,b) and radius r**

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- **For an unknown radius r , unknown gradient direction**






Image space



Hough space

Hough Transform for Circles

- **Circle: center (a,b) and radius r**

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- **For an unknown radius r , *known* gradient direction**





Image space



Hough space

Hough Transform for Circles

For every edge pixel (x,y) :

For each possible radius value r :

For each possible gradient direction θ :

// or use estimated gradient

$$a = x + r \cos(\theta)$$

$$b = y + r \sin(\theta)$$

$$H[a,b,r] += 1$$

end


end

Example: Detecting Circles with Hough


Original



Edges



Votes: Penny




Note: a different Hough Transform (with separate accumulators) was used for each circle radius (quarters vs. penny).


Example: Detecting Circles with Hough

Combined detections


Original



Edges



Votes: Quarter



Coin finding sample images from: Vivek Kwatra

Example: Detecting Circles with Hough



Crosshair indicates results of Hough transform, bounding box found via motion differencing.